



# WALC 2023 Applied AI

## Preventing Overfitting & DL Wrap-Up

---

Prof. Marcelo J. Rovai

[rovai@unifei.edu.br](mailto:rovai@unifei.edu.br)

UNIFEI - Federal University of Itajuba, Brazil

TinyML4D Academic Network Co-Chair



**TINYML4D**

# Preventing Overfitting

# Preventing Overfitting

+Data

+Data

+Data

+Data

# Preventing Overfitting

+Data

+Data

+Data

+Data

What can we do if we  
need more data?

# Preventing Overfitting

+Data

+Data

+Data

+Data

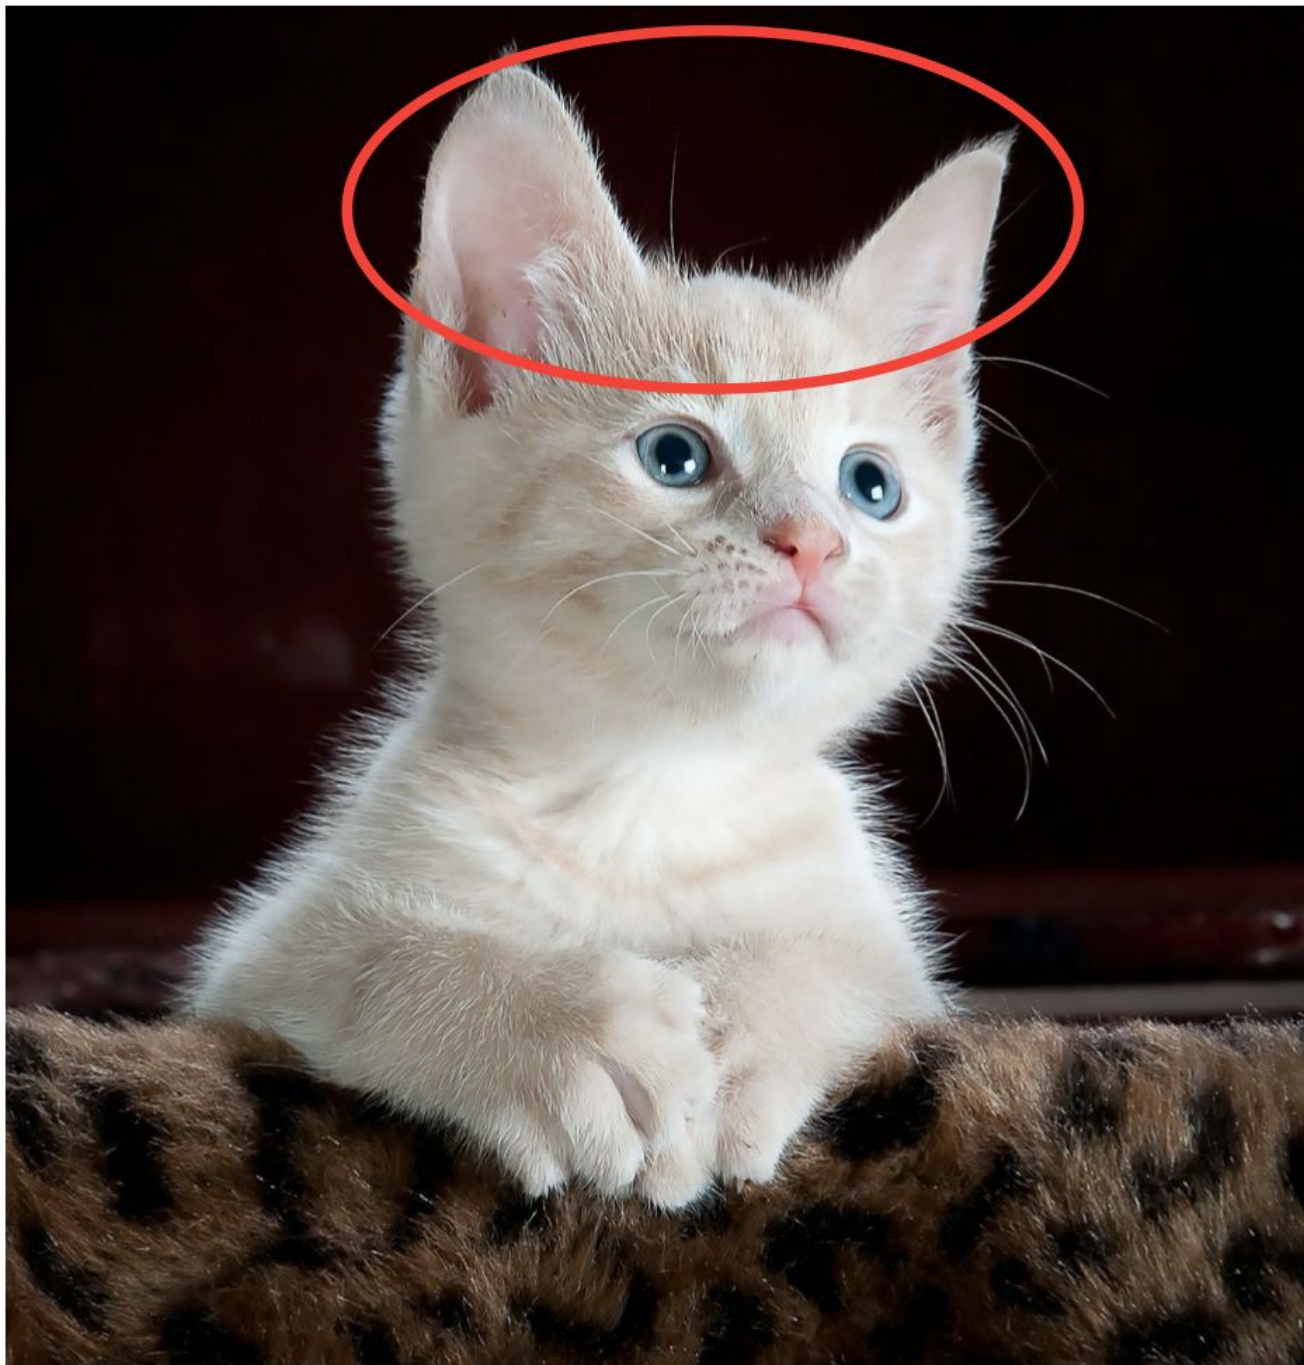
What can we do if we need more data?

- Data Augmentation (artificial)
- Transfer Learning
- Early Stopping
- Dropout Regularization

# Preventing Overfitting

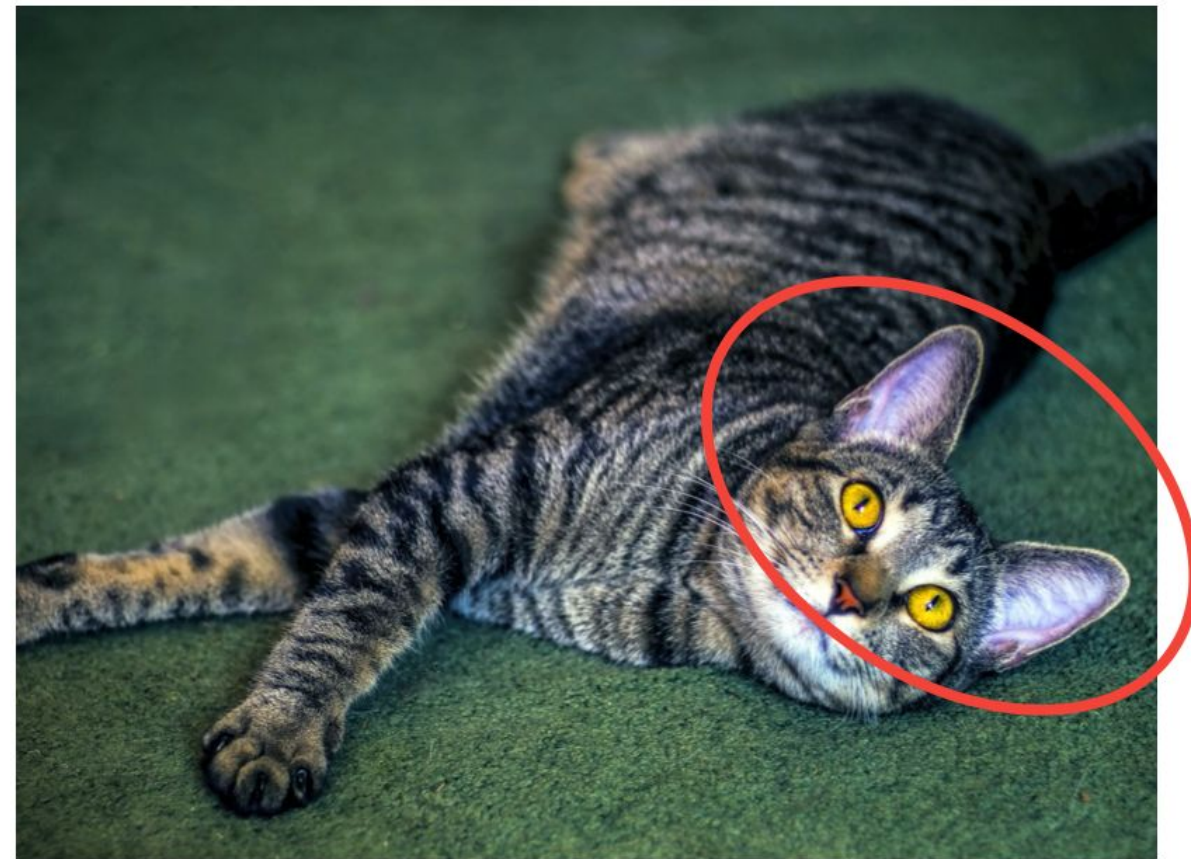
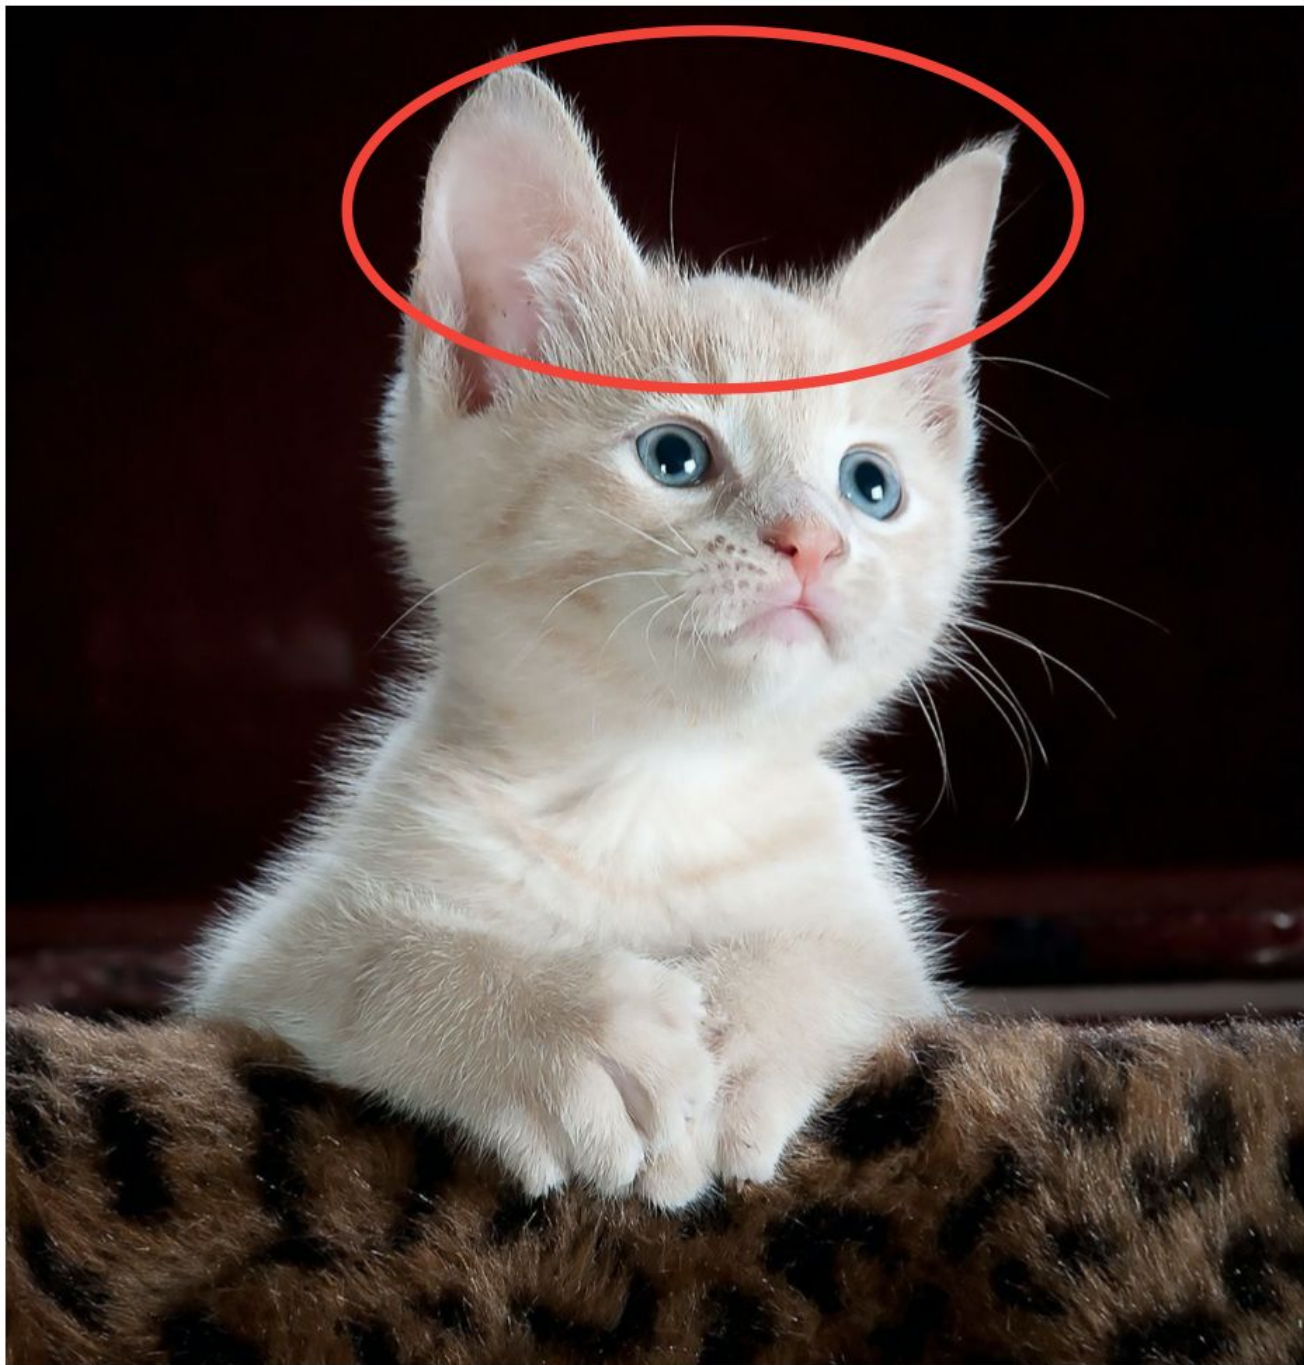
## More Data, Data Augmentation (artificial)

Overfitting generally occurs when there are a small number of training examples. [Data augmentation](#) takes the approach of generating additional training data from your existing examples by augmenting them using random transformations that yield believable-looking images. This helps expose the model to more aspects of the data and generalize better.



← Training

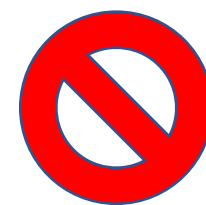





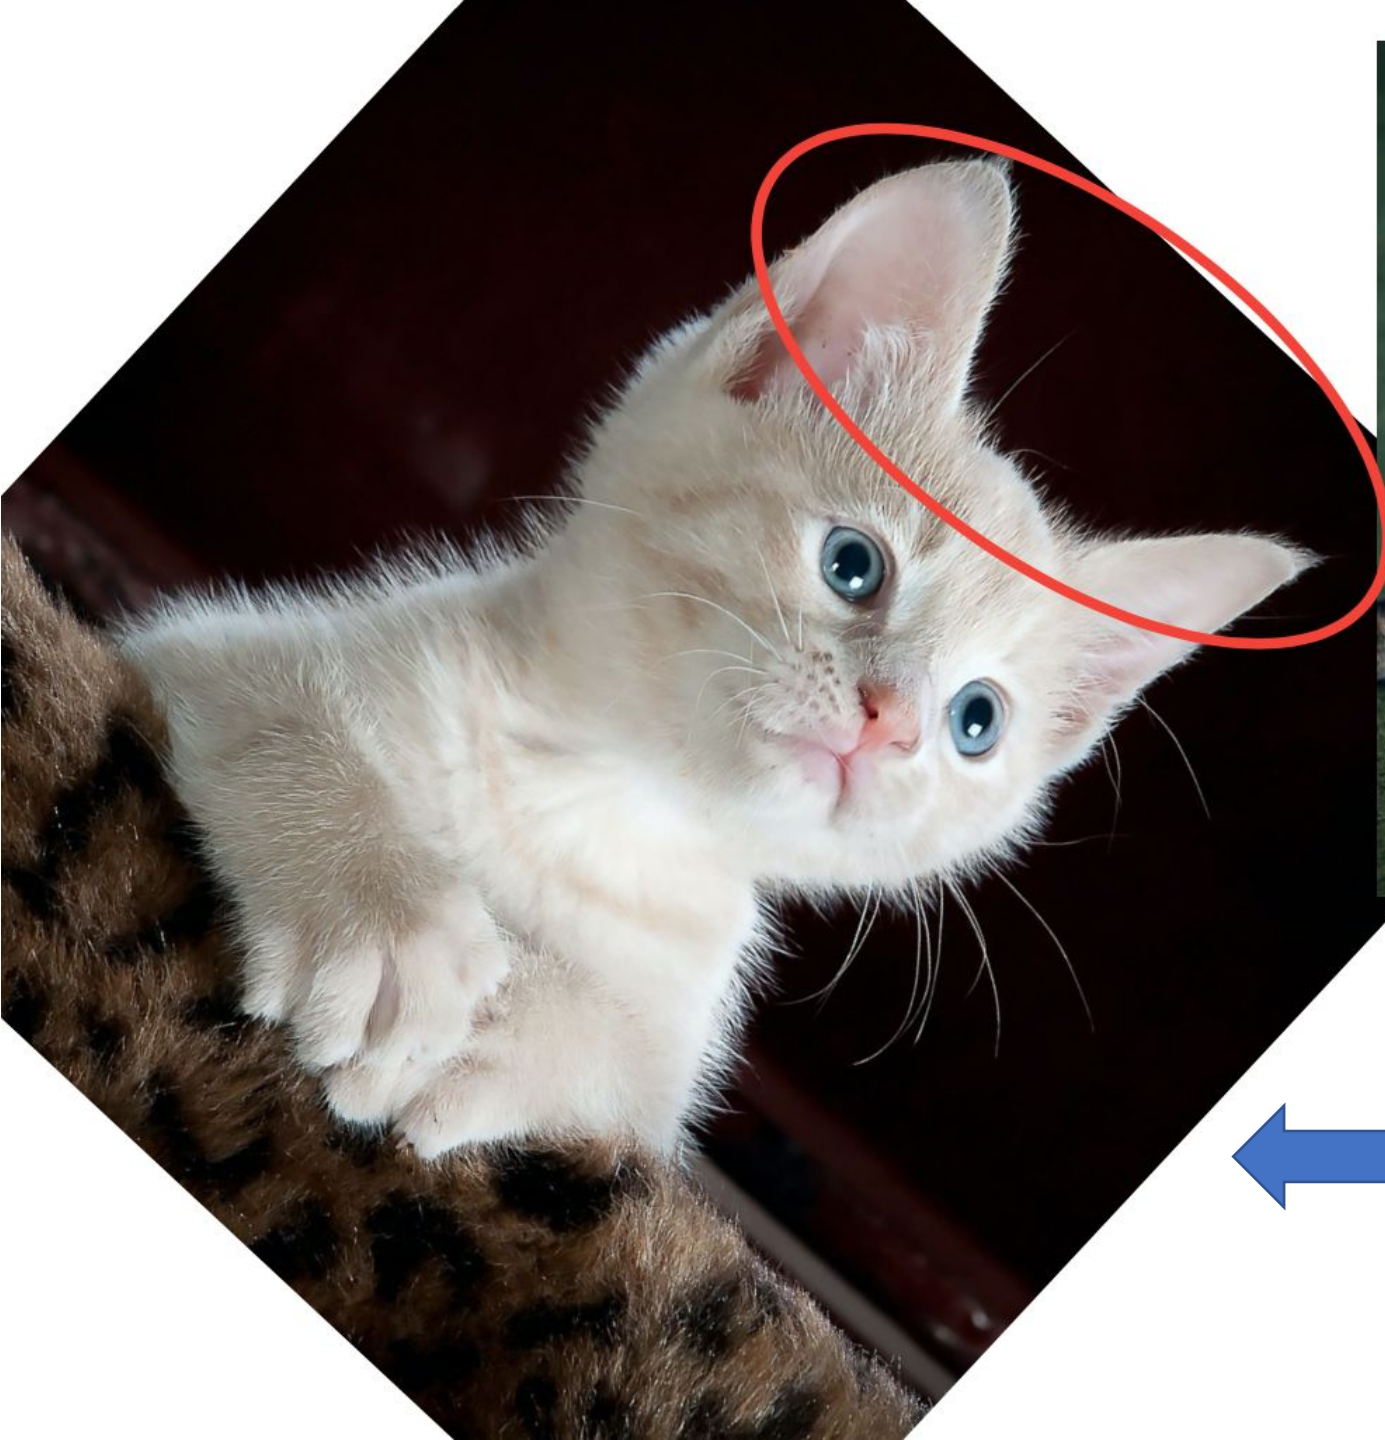
Inference



Training







Training with  
Data Augmentation





# Using Keras preprocessing layers

```
1 data_augmentation = tf.keras.Sequential([  
2     layers.RandomFlip("horizontal_and_vertical"),  
3     layers.RandomRotation(0.2),  
4 ])
```

```
1 plt.figure(figsize=(10, 10))  
2 for i in range(9):  
3     augmented_image = data_augmentation(image)  
4     ax = plt.subplot(3, 3, i + 1)  
5     plt.imshow(augmented_image[0])  
6     plt.axis("off")
```

There are a variety of preprocessing layers you can use for data augmentation including:

- `tf.keras.layers.RandomContrast`,
- `tf.keras.layers.RandomCrop`,
- `tf.keras.layers.RandomZoom`,
- and others.



# Using tf.image

```
1 flipped = tf.image.flip_left_right(image)  
2 visualize(image, flipped)
```

Original image



Augmented image



```
1 rotated = tf.image.rot90(image)  
2 visualize(image, rotated)
```

Original image



Augmented image





# Using tf.image

```
1 saturated = tf.image.adjust_saturation(image, 3)
2 visualize(image, saturated)
```

Original image



Augmented image



```
1 bright = tf.image.adjust_brightness(image, 0.4)
2 visualize(image, bright)
```

Original image



Augmented image



```
1 for i in range(3):
2     seed = (i, 0) # tuple of size (2,)
3     stateless_random_crop = tf.image.stateless_random_crop(
4         image, size=[210, 300, 3], seed=seed)
5     visualize(image, stateless_random_crop)
```

Original image



Augmented image



Original image



Augmented image



Original image



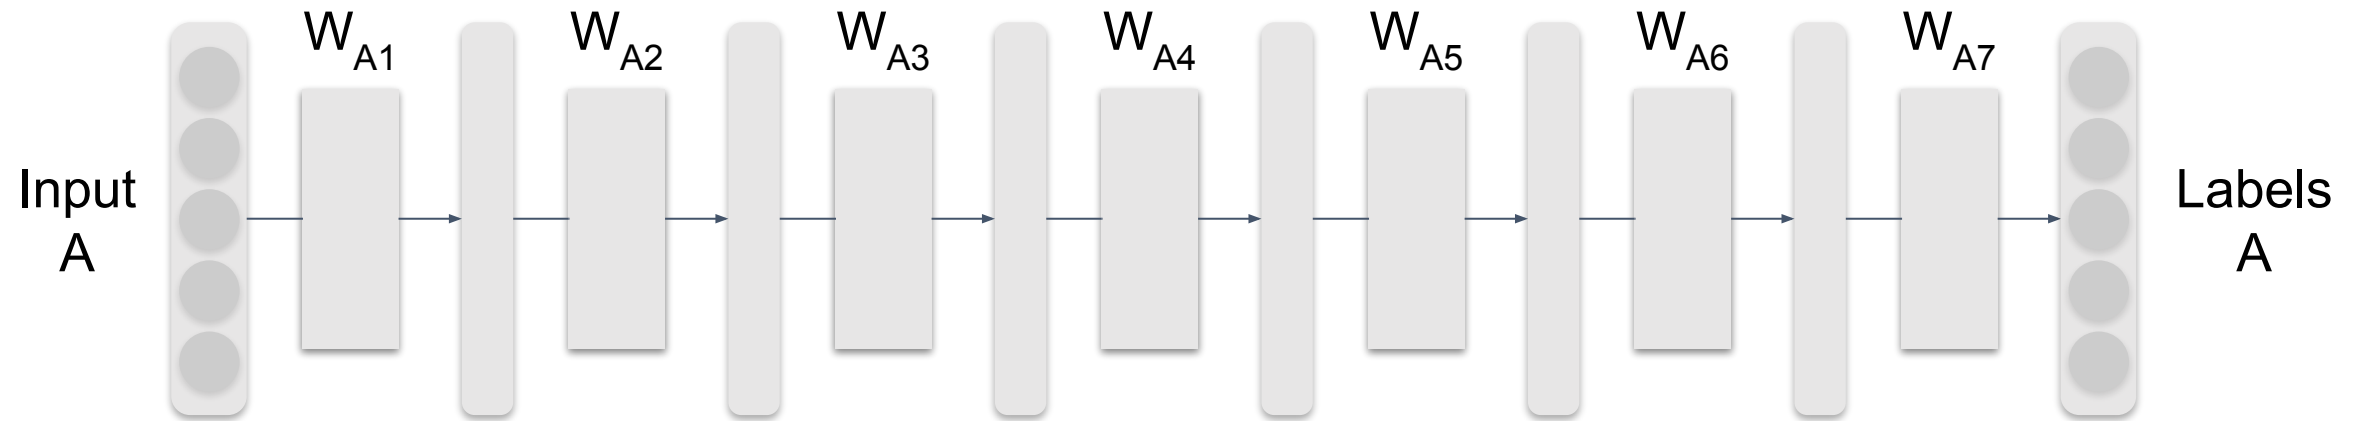
Augmented image



# Preventing Overfitting

## Transfer Learning


# End Result of Training



The end result of the training is to learn the weights of the neural network model.



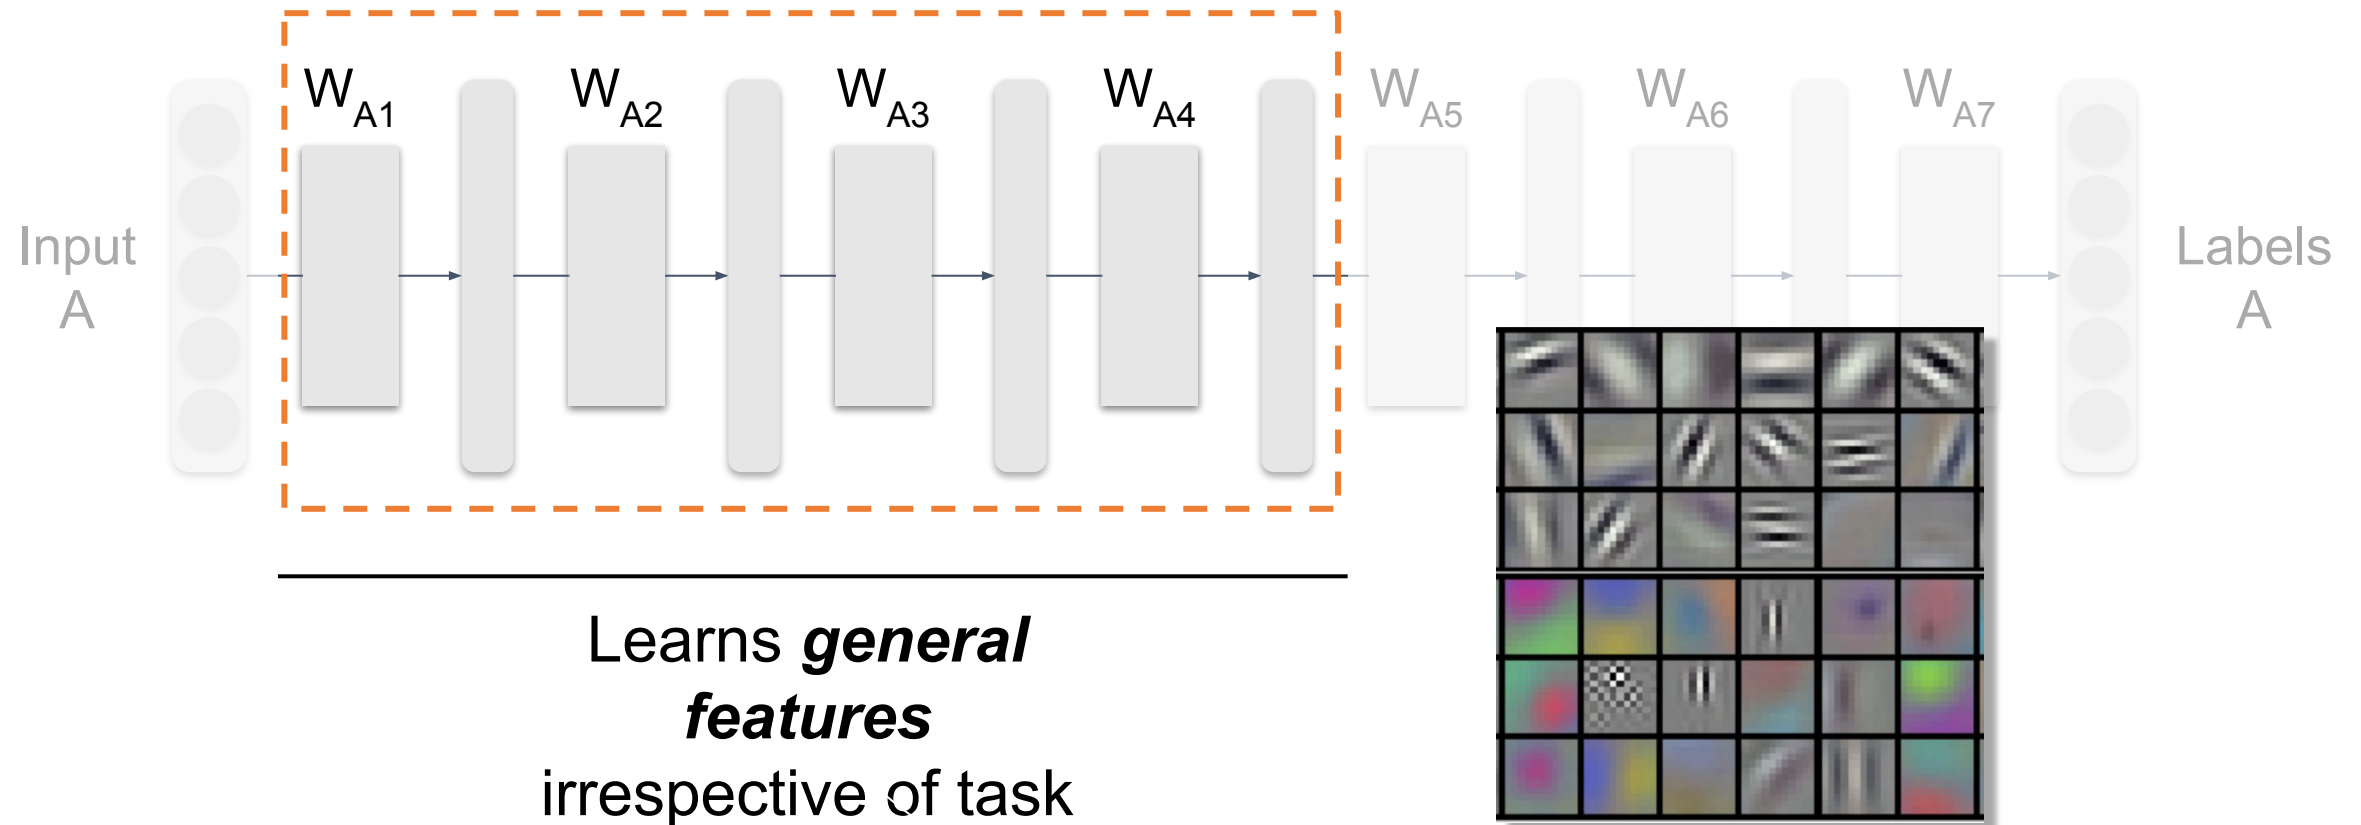
# End Result of Training




---

Learns ***general features***  
irrespective of task

# End Result of Training



# End Result of Training





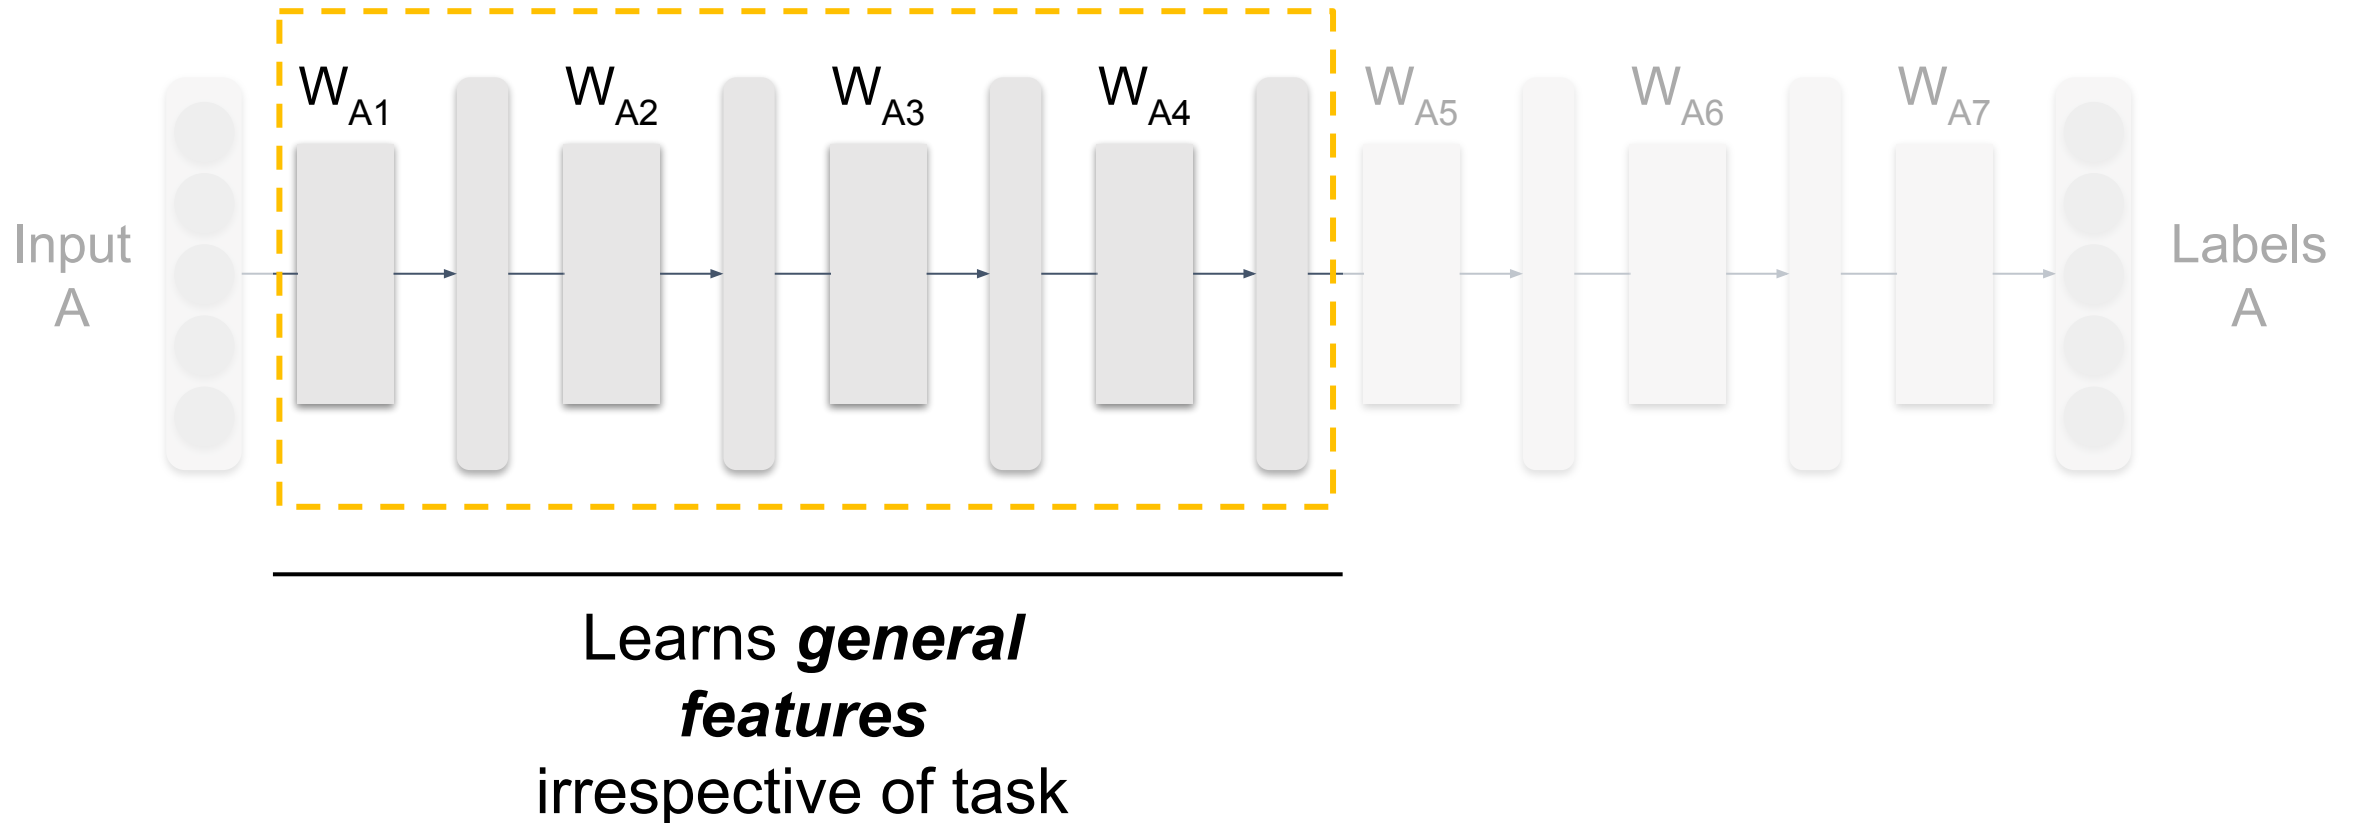


Source: Google




# Transfer Learning

**Reuse** (freeze  
general feature  
extraction)



# Transfer Learning

Train **only** last  
few layers

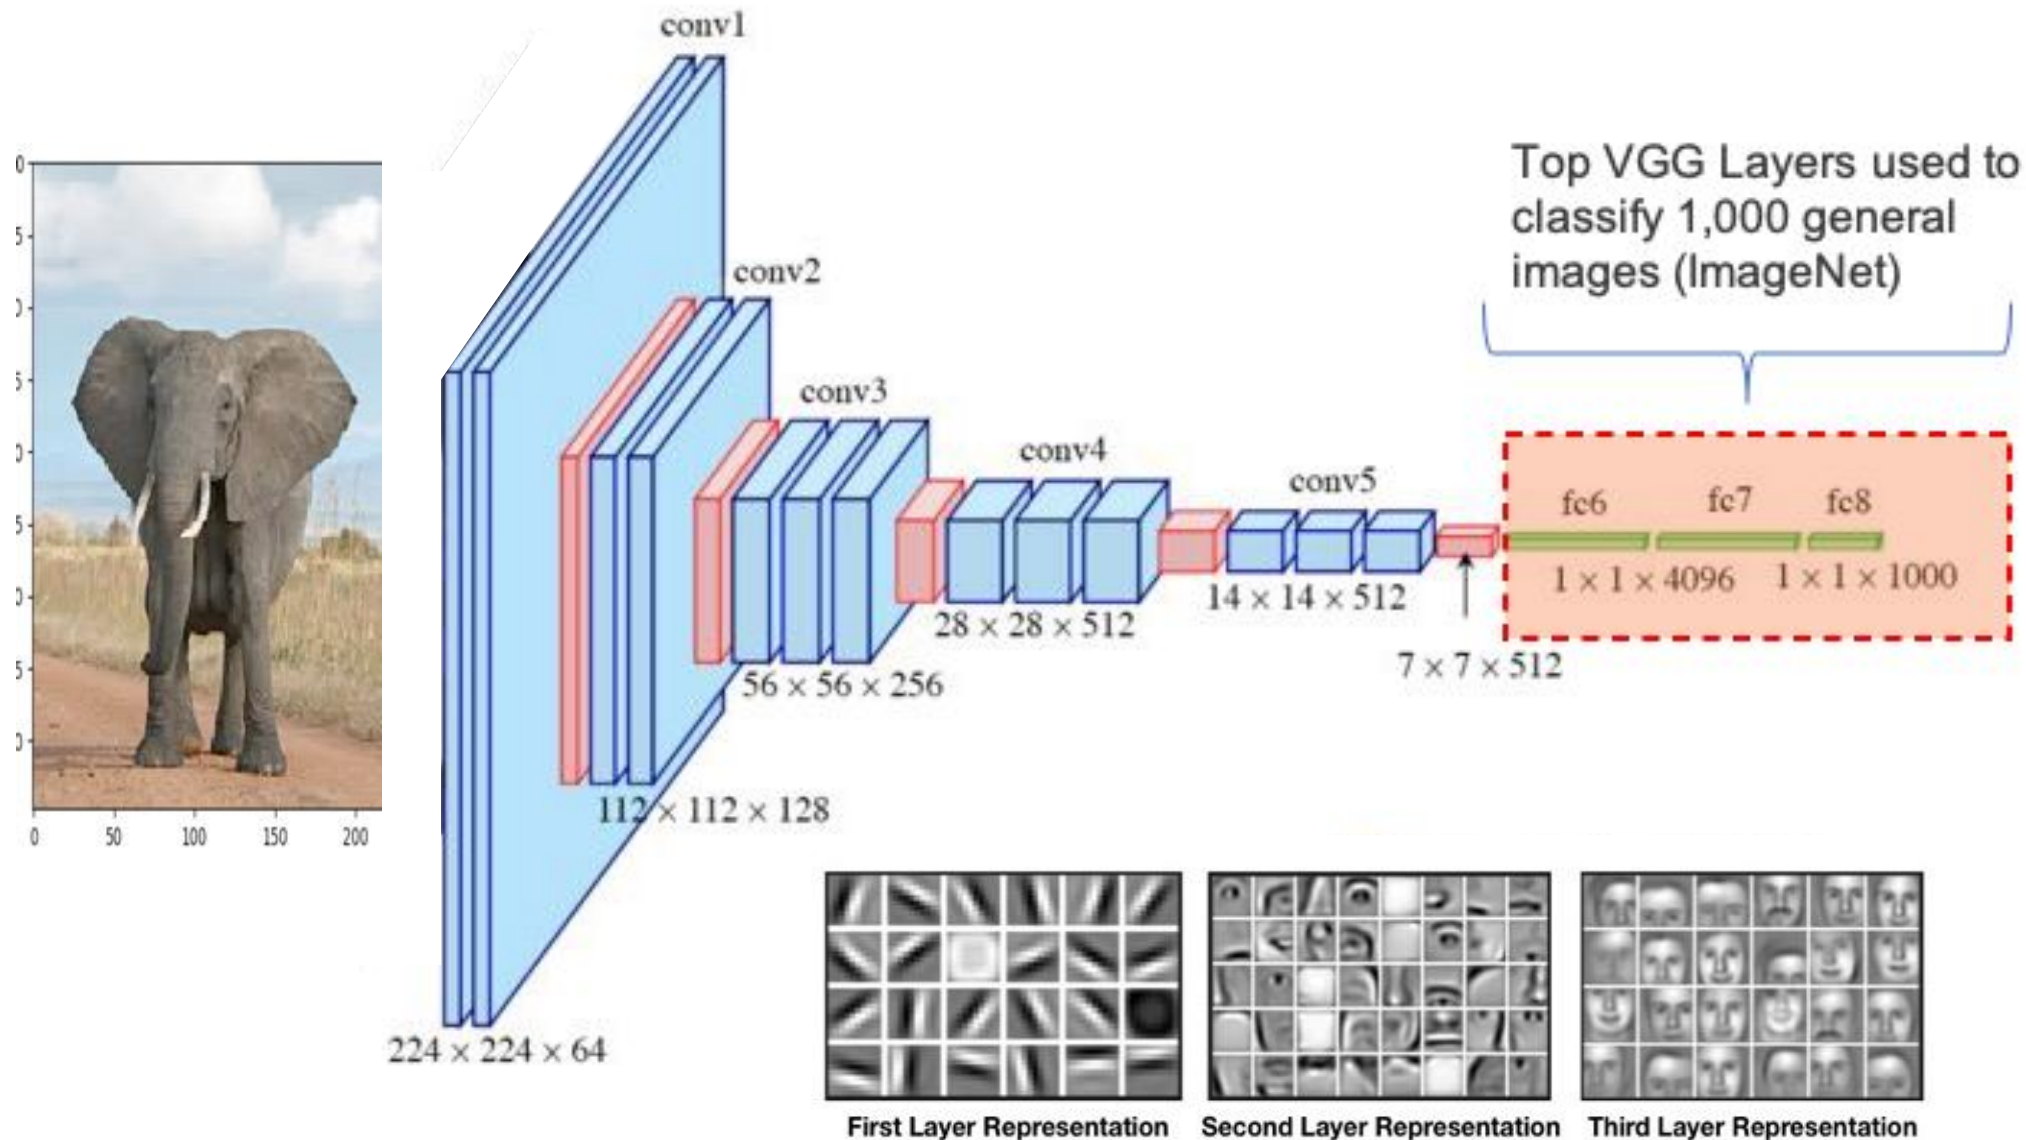




covidXray

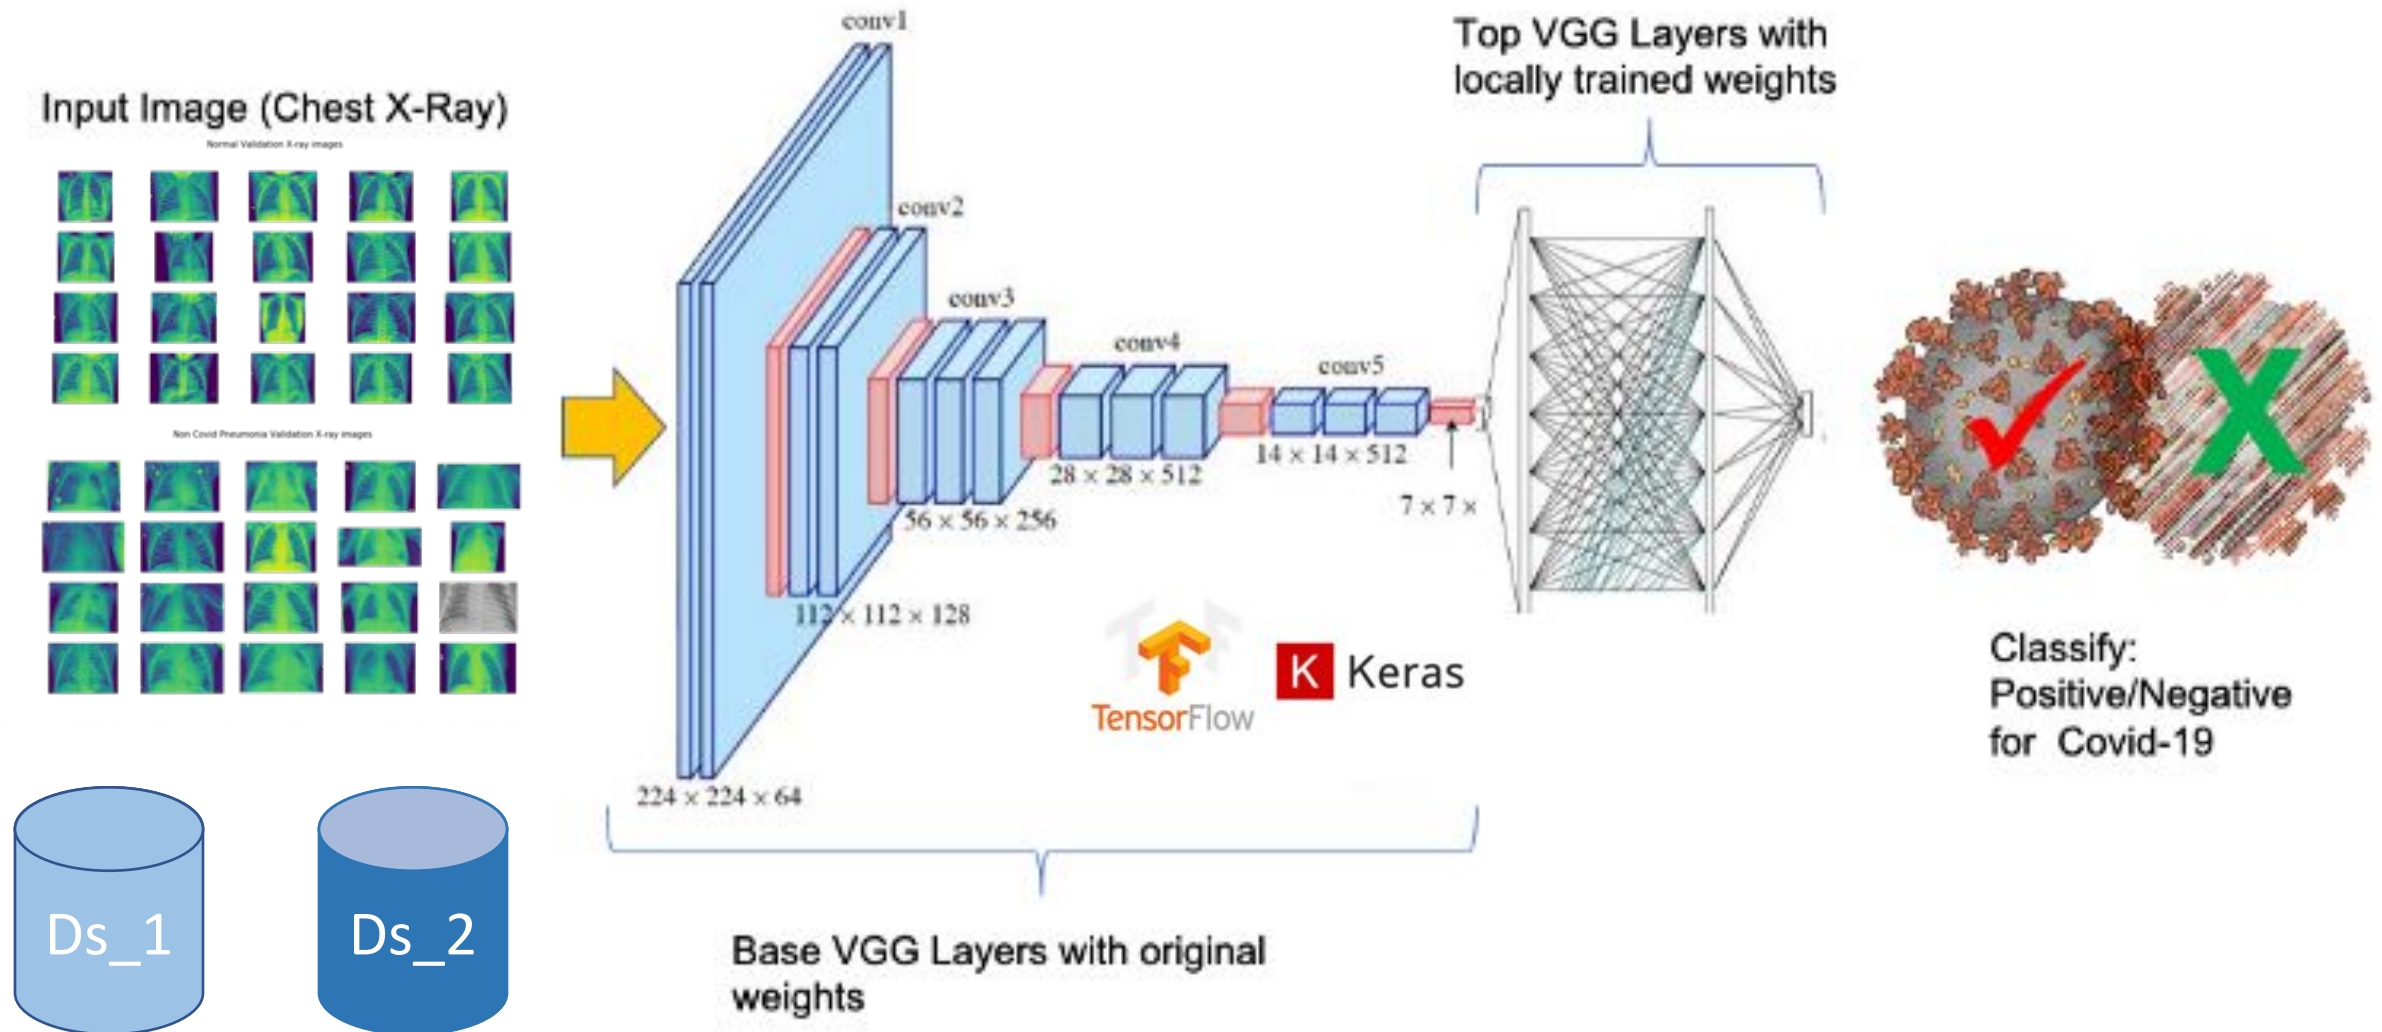
Detecting Covid-19 in Chest X-Ray images

# VGG-16 Convolutional Neural Network Model

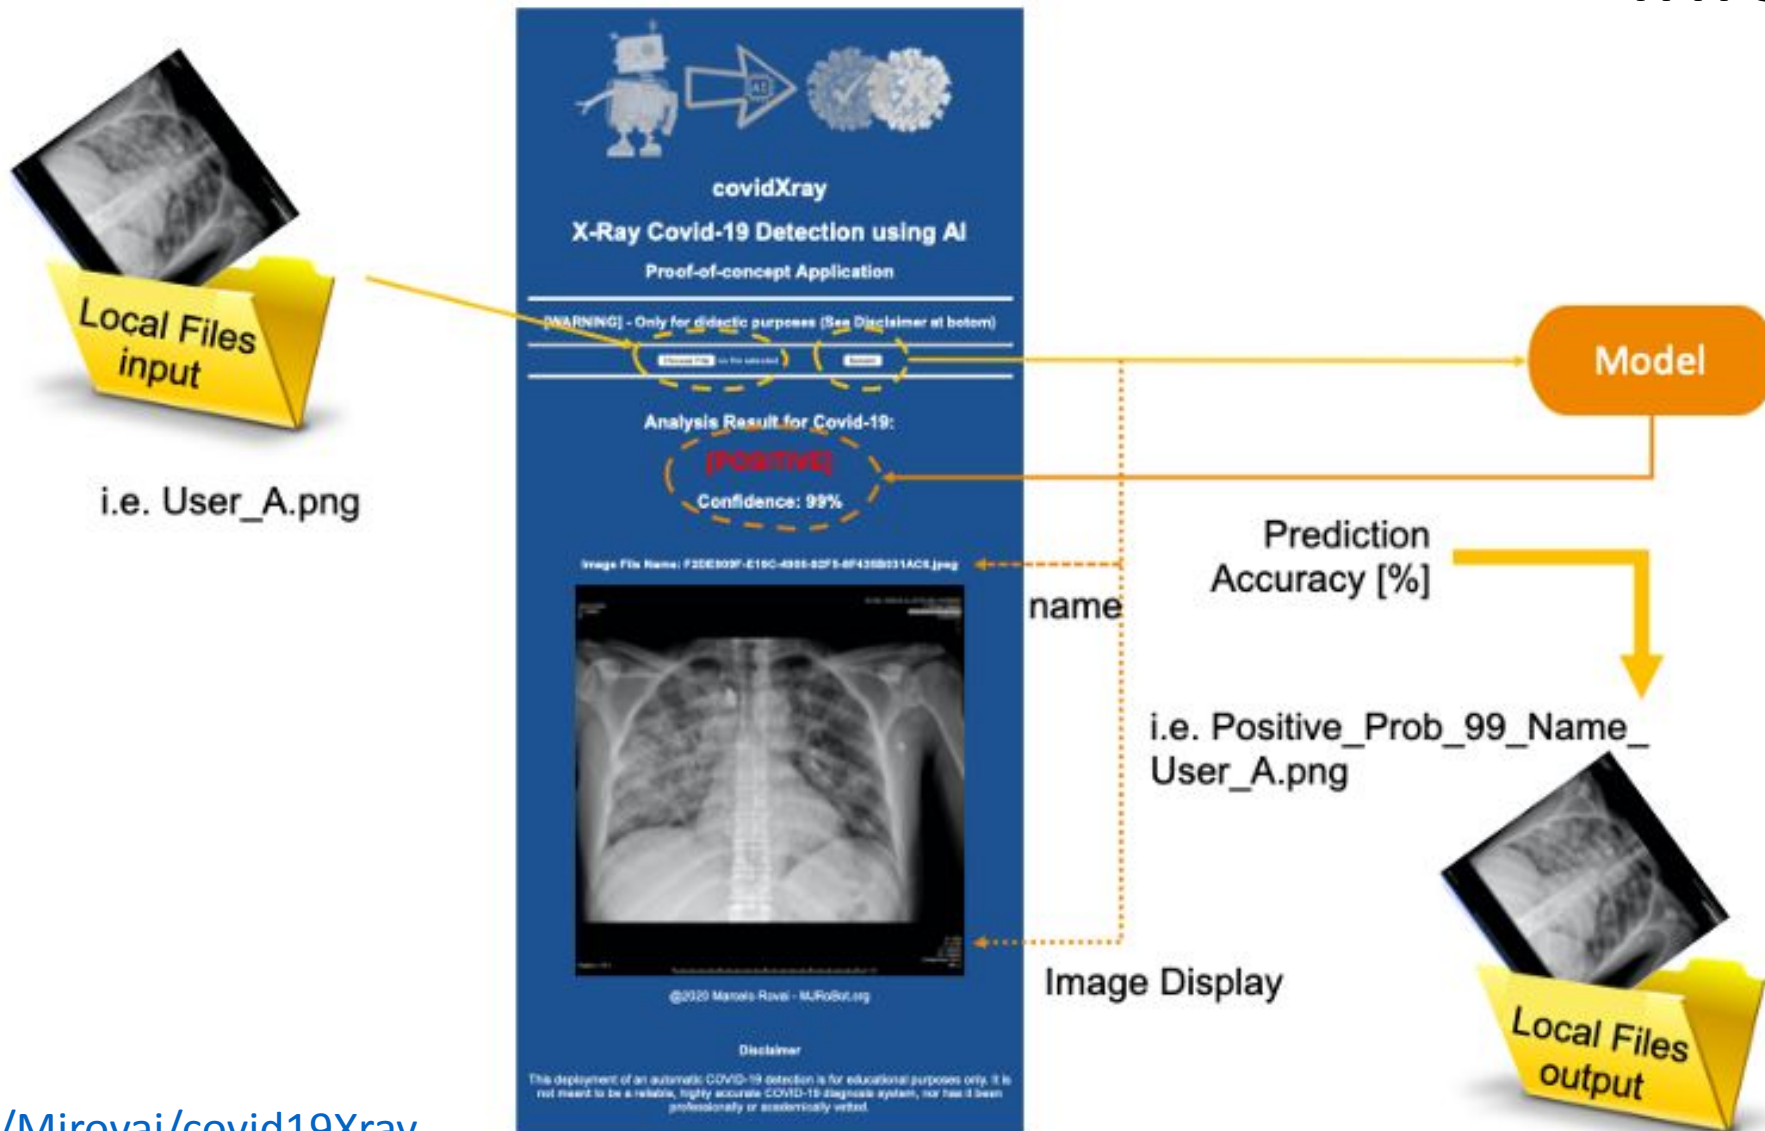


**Elephant**  
**93.0%**

# Training the model (Transfer Learning)



# Inference



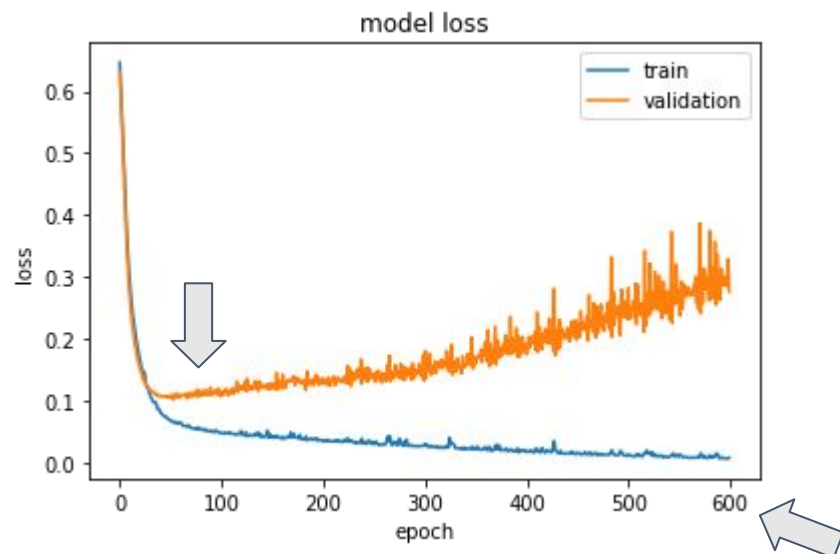
<https://github.com/Mjrovai/covid19Xray>

# Preventing Overfitting

Early Stopping & Dropout Regularization

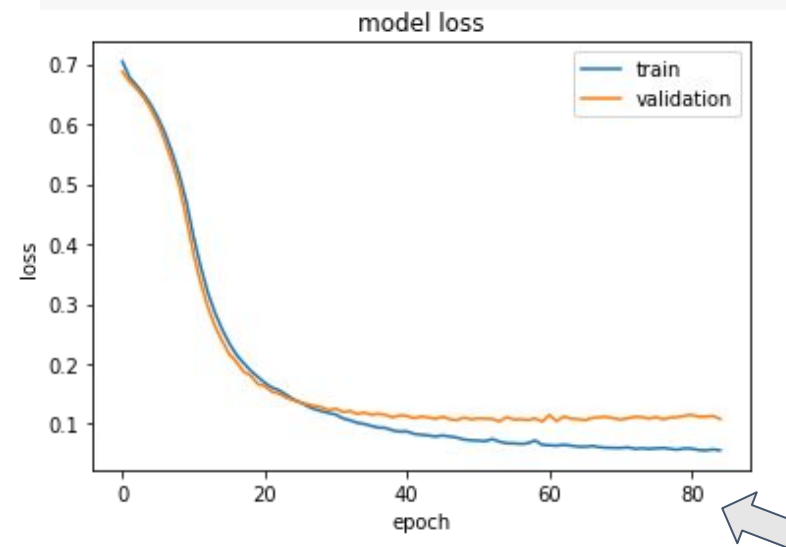
# Early Stopping

```
history = model.fit(X_train,  
                    y_train,  
                    epochs=600,  
                    validation_data=(X_test, y_test),  
                    verbose=1  
                    )
```



```
from tensorflow.keras.callbacks import EarlyStopping  
  
early_stop = EarlyStopping(monitor='val_loss',  
                           mode='min',  
                           verbose=1,  
                           patience=25)
```

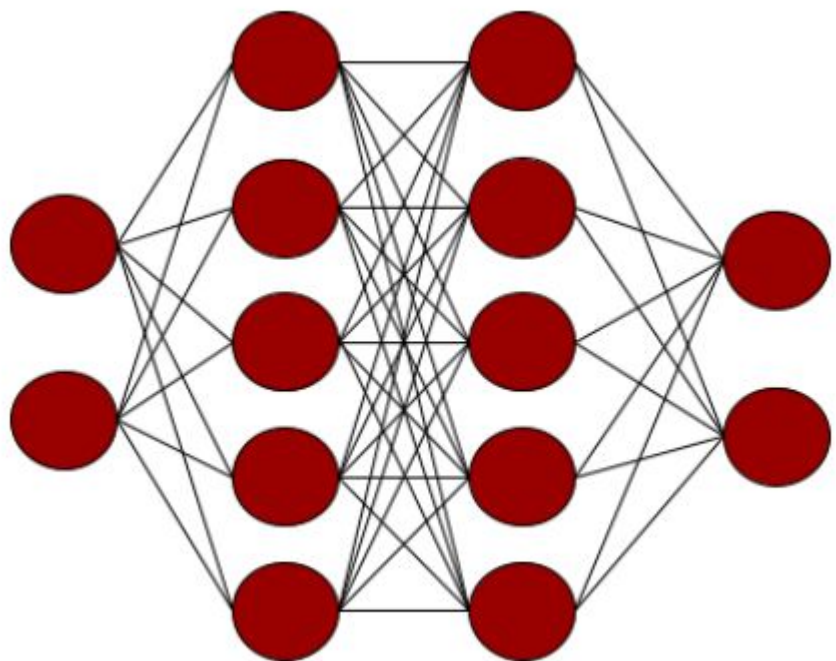
```
history = model.fit(x=X_train,  
                    y=y_train,  
                    epochs=600,  
                    validation_data=(X_test, y_test),  
                    verbose=1,  
                    callbacks=[early_stop]  
                    )
```





# Dropout Regularization

```
model = tf.keras.models.Sequential([  
    tf.keras.layers.Flatten(input_shape=(28, 28)),  
    tf.keras.layers.Dense(256, activation=tf.nn.relu),  
    tf.keras.layers.Dense(128, activation=tf.nn.relu),  
    tf.keras.layers.Dense(64, activation=tf.nn.relu),  
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)])
```

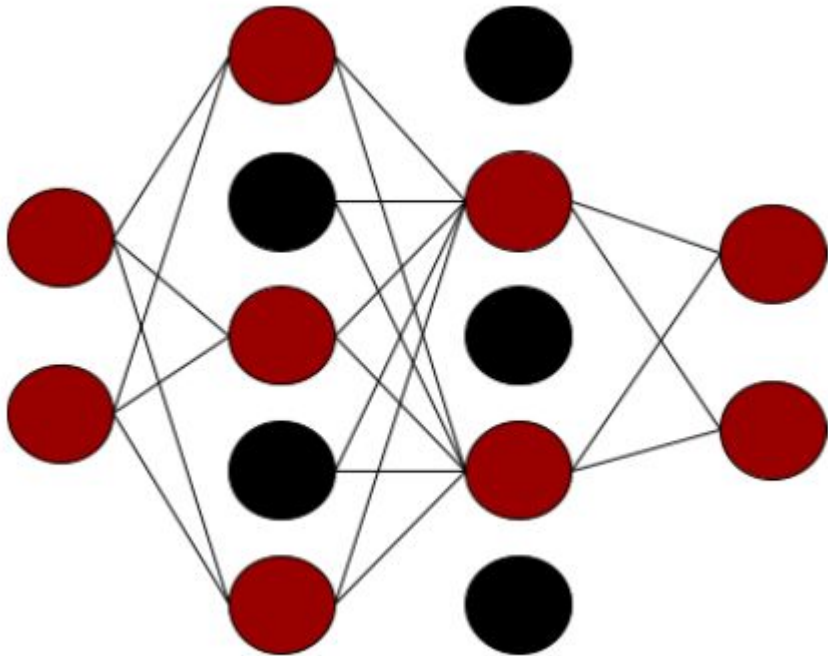


## Fashion MNIST Dataset

- 20 Epochs
- 94.0% Accuracy on Train Data
- 88.5% Accuracy on Validation Data

# Dropout Regularization

```
model = tf.keras.models.Sequential([  
    tf.keras.layers.Flatten(input_shape=(28,28)),  
    tf.keras.layers.Dense(256, activation=tf.nn.relu),  
    tf.keras.layers.Dropout(0.2),  
    tf.keras.layers.Dense(128, activation=tf.nn.relu),  
    tf.keras.layers.Dropout(0.2),  
    tf.keras.layers.Dense(64, activation=tf.nn.relu),  
    tf.keras.layers.Dropout(0.2),  
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)])
```



## Fashion MNIST Dataset

- 20 Epochs
- 89.5% Accuracy on Train Data
- 88.3% Accuracy on Validation Data



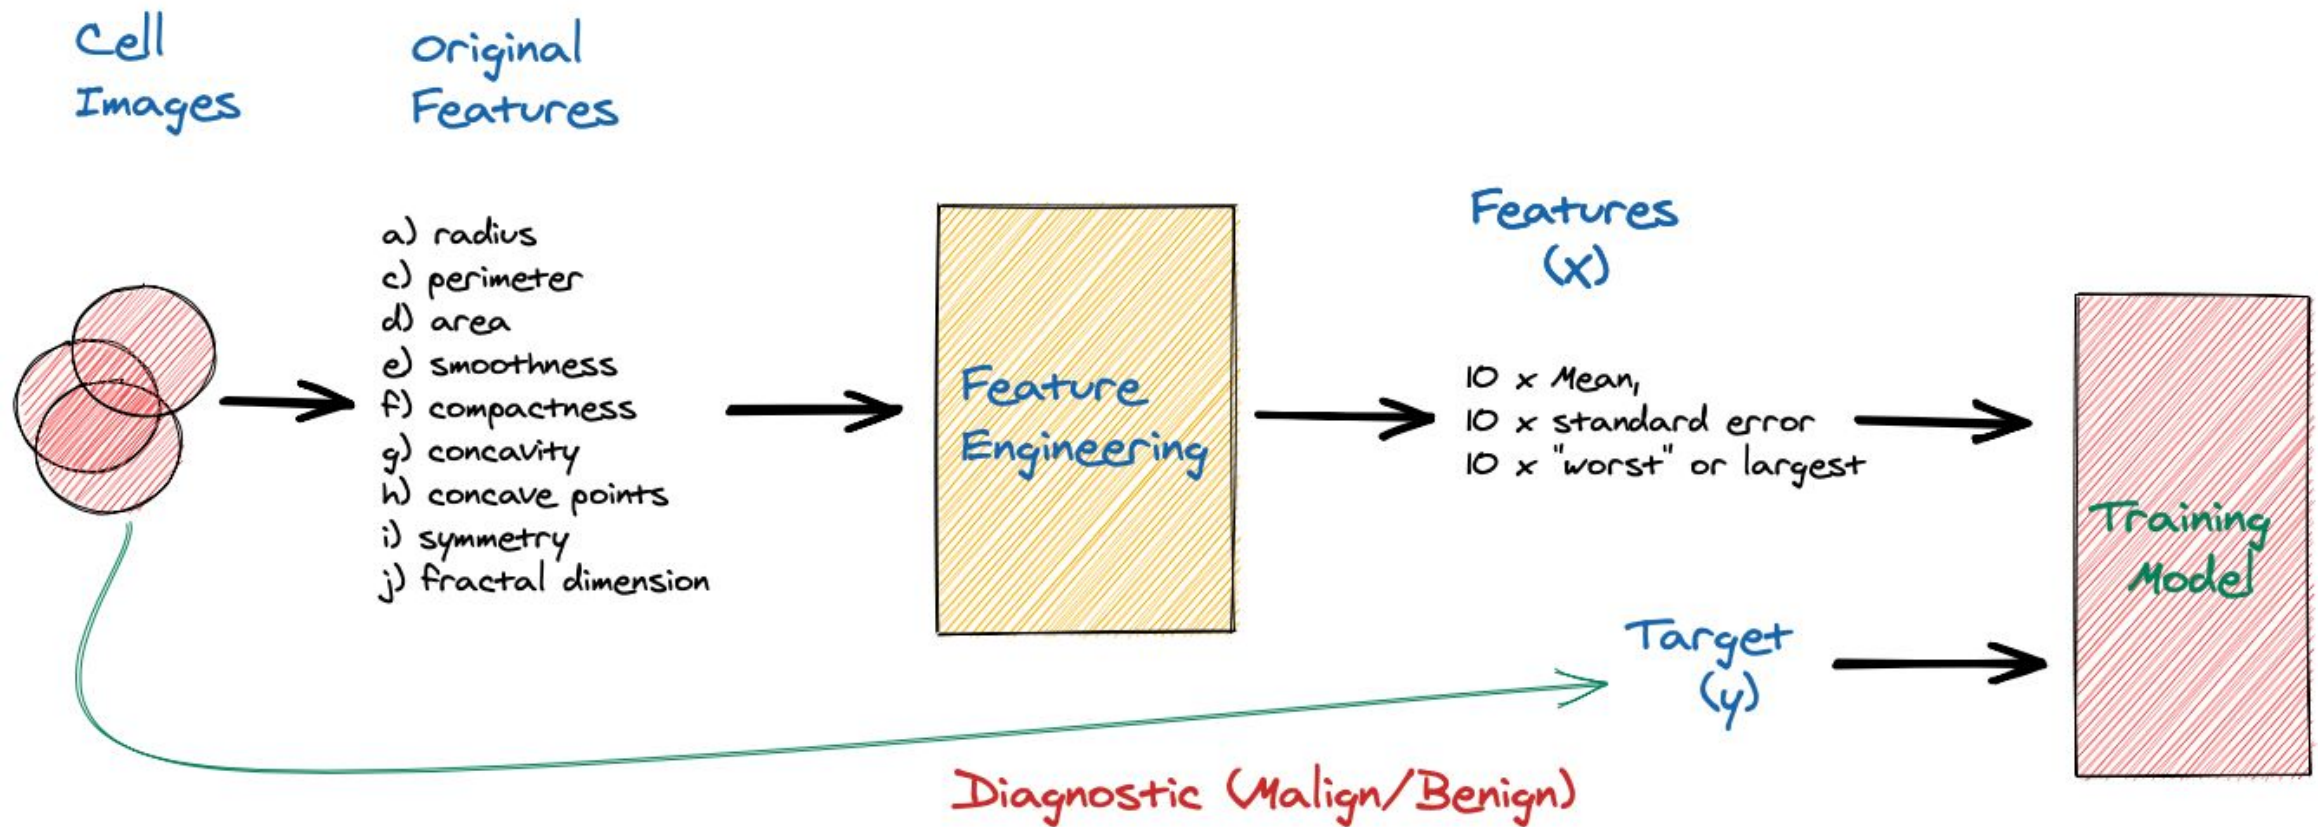
Removing a random number of neurons and connections (in this example, 20%), reduces the chances of the neurons becoming overspecialized and the model will generalize better, reducing the overfit.

# Wisconsin Diagnostic Breast Cancer (WDBC)

## Optional Homework

[Breast\\_Cancer\\_Classification.ipynb](#)



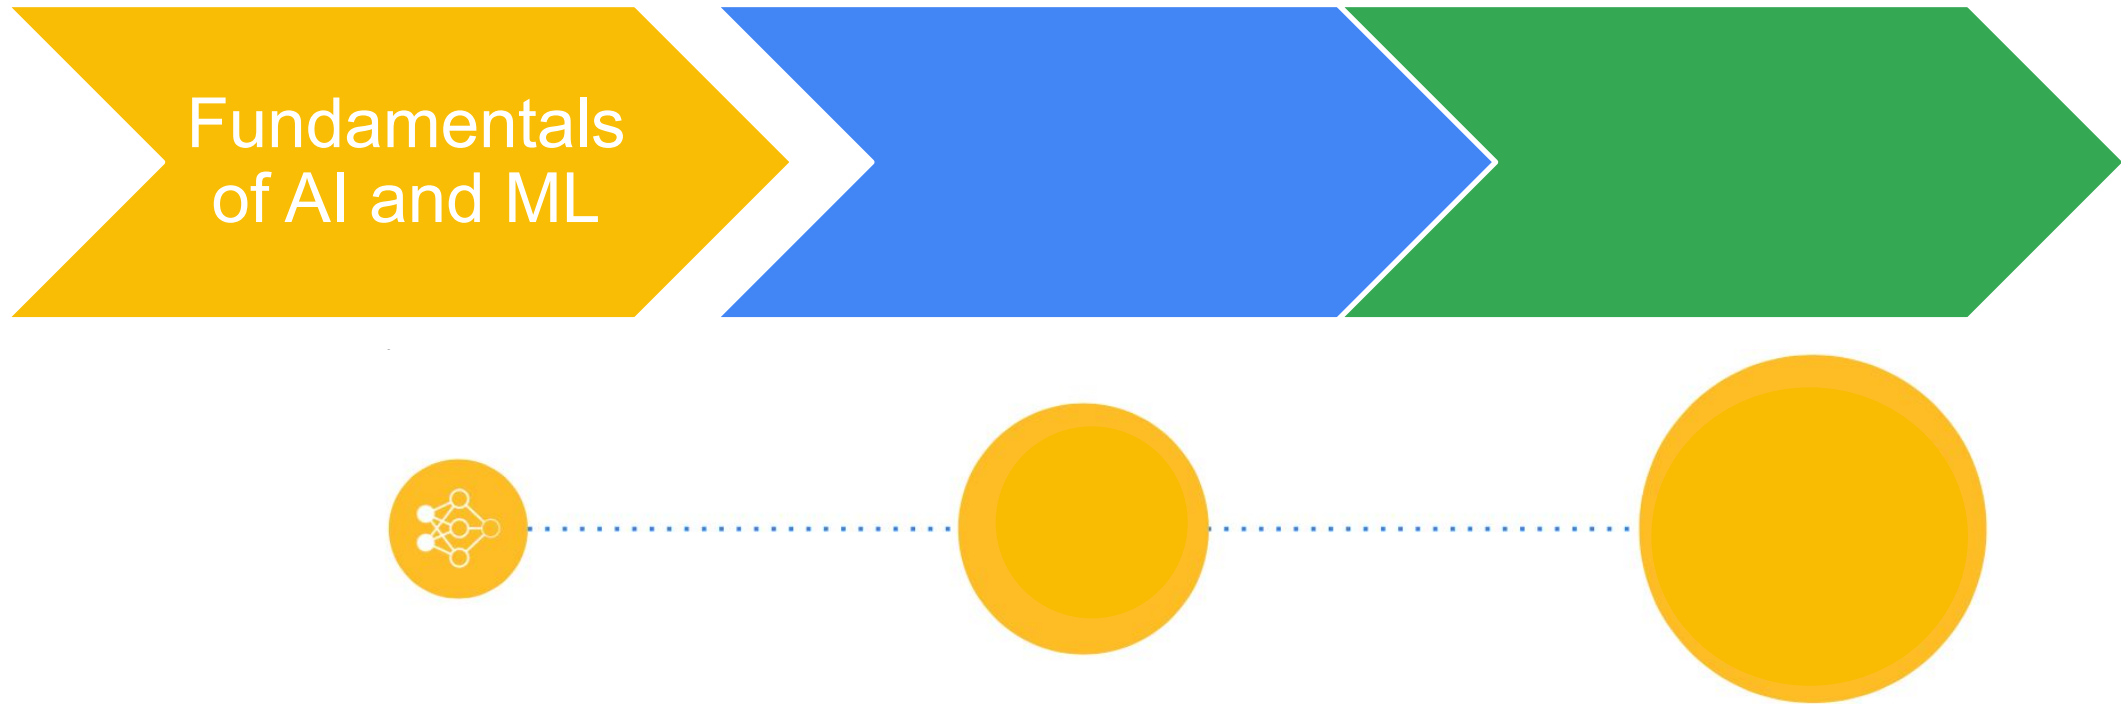


UCI ML Breast Cancer Wisconsin (Diagnostic) datasets. <https://goo.gl/U2Uwz2>

# Deep Learning Wrap-Up

# What have we learned so far?

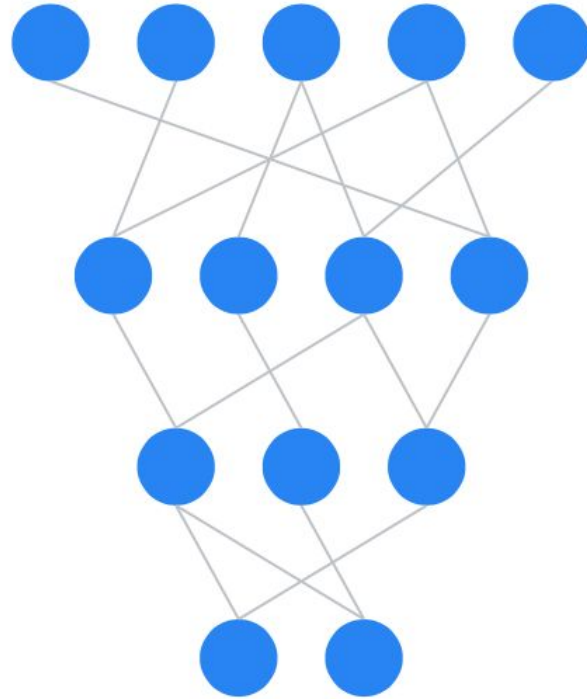
## Part 1



In Part 1, while discussing what is the **language of machine learning**, we introduced ML with TensorFlow.



# Total Recall from **Part 1**



# “Language” for Part 1

Neural Network

Gradient Descent

Loss Function

Training Data

Training

Validation Data

Inference

Test Data

Features

Classification

Filters

Overfitting

Regression

Kernels

Data augmentation

Responsible AI

CNNs

DNNs

Preprocessing

# “Language” for Part 1

Training Data

Neural Network

**Training**

Validation Data

Gradient Descent

**Inference**

Test Data

Loss Function

**Features**

Classification

**Filters**

**Overfitting**

Regression

**Kernels**

Data augmentation

Responsible AI

**CNNs**

**DNNs**

Preprocessing

# “Language” for Part 1

**Training Data**

Neural Network

Training

**Validation Data**

Gradient Descent

Inference

**Test Data**

Loss Function

Features

Classification

Filters

Overfitting

Kernels

Regression

**Data augmentation**

CNNs

DNNs

Responsible AI

**Preprocessing**



# “Language” for Part 1

Training Data

Neural Network

Training

Validation Data

Gradient Descent

Inference

Test Data

Loss Function

Features

**Classification**

Filters

Overfitting

Kernels

**Regression**

CNNs


Data augmentation

DNNs

**Responsible AI**

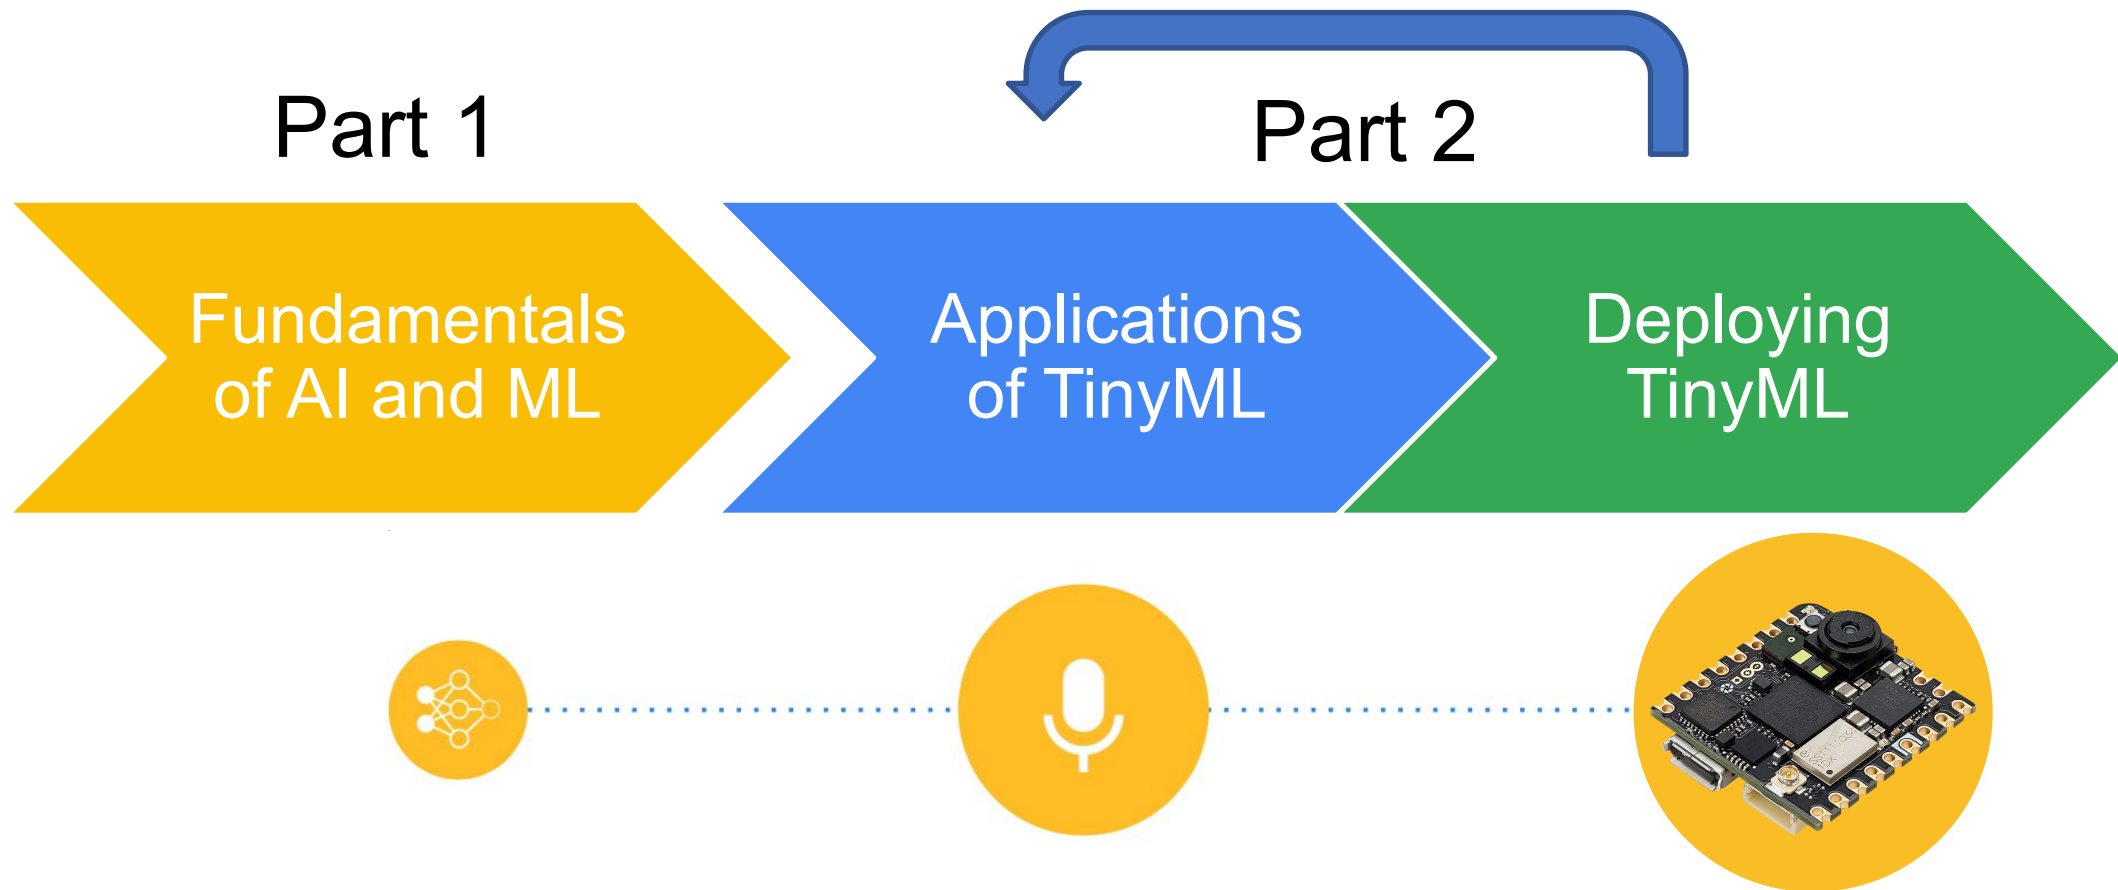
Preprocessing

# What will we learn?





In Part 2, we will get a sneak peek into the variety of different TinyML applications, as keyword spotting (“Alexa”), gesture recognition, understand how to leverage the sensors, and so forth.

# What will we learn?



In Part 2, we will **also** learn how to deploy models on a real microcontroller. Along the way we will explore the challenges unique to and amplified by TinyML (e.g., preprocessing, post-processing, dealing with resource constraints).








Train a model

Convert  
model

Optimize  
model

Deploy  
model at  
Edge

Make  
inferences  
at Edge





# Thanks



TINYML4D

